

Adaptive Precision Attention: Evolutionary Discovery to Hardware-Verified Ratio Gates

Chaithu Talasila

<https://github.com/LonghornSilicon/attention-compute-unit>

Abstract

We present an evolutionary approach to discovering per-block precision allocation policies for mixed-precision attention, then trace a complete path from policy discovery to hardware-verified implementation. Using OpenEvolve, an LLM-guided evolutionary search framework, we evolve a precision policy over 14 per-block statistics and 6 precision levels. The search converges on a two-threshold policy using softmax entropy and outlier presence. We then show that entropy—which requires 32,000 transcendental operations per tile—is replaceable by a pre-softmax ratio signal ($\max |S| / \text{mean} |S|$) with 100% decision agreement and zero dangerous misses at threshold 10. The hardware implementation requires no division: $\max \times N > 10 \times \text{sum}$, a left shift, two shifts-and-add for multiply-by-10, and one comparator.

Validated on real Qwen2/Qwen2.5 language models at 512-token sequences, 91–97% of attention tiles are INT8-safe across model sizes from 0.5B to 3B parameters (24–36 layers). INT8 safety *improves* with scale: the 1.5B model reaches 97.1% INT8-safe tiles. Fixed-point simulation confirms 100% decision agreement at every tested bit width (4, 6, 8, 10, 12, 16 bits), with a 16-bit sum accumulator and 20-bit comparator sufficient. The precision controller occupies a handful of logic gates and 64 bytes of per-layer threshold registers—directly implementable as an ASIC attention accelerator.

1 Introduction

The attention mechanism is the computational core of transformer-based language models, and its efficient implementation has been a sustained focus of systems research. FlashAttention [Dao et al., 2022] and its successors [Dao, 2024, Shah et al., 2024] achieve near-peak hardware utilization through tiled, fused com-

putation that avoids materializing the full $N \times N$ attention matrix. FlashAttention-3 introduces FP8 quantization on Hopper GPUs, achieving up to 1.2 PFLOPs/s, but treats precision as a global choice: all blocks are computed at the same numerical format.

In practice, different blocks of the attention matrix have different precision requirements. Blocks where the softmax distribution is peaked and input values contain statistical outliers are sensitive to quantization error. Blocks where attention is uniformly distributed are robust to aggressive compression. This motivates *adaptive precision attention*, where each block is quantized to the minimum precision that preserves output quality.

The challenge is determining which blocks need protection. We pose precision allocation as a program synthesis problem and solve it with evolutionary search over per-block statistics. We then carry the discovered policy all the way to fixed-point hardware simulation, validating that the signal survives real LLM activations and integer arithmetic.

Contributions.

- An evolutionary framework for precision policy discovery over 14 per-block statistics and 6 precision levels, converging on two comparisons and one AND gate.
- Proof that softmax entropy (32K transcendental ops/tile) is replaceable by a pre-softmax ratio signal using only a max tree, an adder tree, and one comparator—with 100% agreement and zero dangerous misses.
- Real LLM validation on Qwen2-0.5B, Qwen2-1.5B, Qwen2.5-3B, and Phi-2 (2.7B): 91–97% of tiles are INT8-safe; INT8 coverage *improves* with model scale.
- Fixed-point simulation confirming 100% agreement at 4-bit score representation; hardware reg-

ister budget of 16-bit sum + 20-bit comparator.

- A working Triton kernel prototype (12/12 correctness tests passing) and KV-cache module with uniform $2\times$ compression.

2 Related Work

Efficient attention. FlashAttention [Dao et al., 2022] introduced tiled, IO-aware attention computation. FlashAttention-2 [Dao, 2024] improved parallelism. FlashAttention-3 [Shah et al., 2024] targets Hopper GPUs with warp specialization and uniform FP8 block quantization. All three treat precision as a global setting.

Adaptive quantization. Boroujeni et al. [2026] propose per-token KV-cache quantization with four precision levels selected by entropy and attention variance, achieving 17.75% latency reduction. SmoothQuant [Xiao et al., 2023] redistributes quantization difficulty between activations and weights. GPTQ [Frantar et al., 2023] uses second-order information for weight quantization. Our work focuses on attention computation precision using evolutionary search, targeting ASIC implementation.

Program synthesis. OpenEvolve [CodeLion, 2025] extends FunSearch [Romera-Paredes et al., 2024] with multi-objective optimization and island-based population management. We use it as the search engine with Claude Sonnet as the LLM backbone.

3 Method

3.1 Problem Formulation

We define a precision policy as a function mapping per-block statistics to a precision level:

$$\text{policy} : \mathcal{S} \rightarrow \{\text{fp16}, \text{fp8_e4m3}, \text{fp8_e5m2}, \text{int8}, \text{fp4}, \text{int4}\}$$

For each query block $\mathbf{Q}_i$ and key-value block $(\mathbf{K}_j, \mathbf{V}_j)$, we compute block statistics, query the policy, and apply a quantize-dequantize roundtrip before accumulation (always in FP32).

3.2 Block Statistics

For each block pair (i, j) we extract 14 statistics: L2 norms of $\mathbf{Q}$, $\mathbf{K}$, $\mathbf{V}$ blocks; statistics of pre-softmax

scores $\mathbf{S} = \mathbf{QK}^\top / \sqrt{d}$; the Shannon entropy of the row-wise softmax; causal distance; block indices; and an outlier flag:

$$\text{has_outlier} = \max |\mathbf{S}| > \mu_{|\mathbf{S}|} + 10 \sigma_{|\mathbf{S}|}$$

Block size is 128 tokens, matching FlashAttention-3’s tile size.

3.3 Evolutionary Search

We use OpenEvolve [CodeLion, 2025] with 50 candidates across 3 islands, 80 iterations of diff-based mutation, Claude Sonnet (80%) and Haiku (20%) as LLM backbone, and a fitness function weighting 50% accuracy $(1/(1 + \text{RMSE} \cdot 10))$ vs FP64) + 30% compression + 20% reliability.

3.4 Evaluation Workloads

We evaluate on 12 workloads spanning $N \in \{512, 2048, 4096, 8192\}$, $d \in \{64, 128\}$, with and without causal masking, and with 0.1% of elements scaled by $100\times$ to simulate LLM activation outliers.

4 Results

4.1 Converged Policy

After 80 iterations (combined score 0.745, target >0.65), the search converged on:

```
def precision_policy(block_stats):
    has_outlier = block_stats.get('
        has_outlier', False)
    entropy = block_stats.get('entropy',
        4.3)
    if has_outlier and entropy < 2.0:
        return "fp16"
    return "int8"
```

Listing 1: Evolved precision policy (2 of 14 features used).

All intermediate formats (FP8, FP4, INT4) were eliminated by the search.

4.2 Accuracy vs. Baselines

Table 1 summarises aggregate metrics. Table 2 gives per-workload detail. Figure 1 shows the log-scale RMSE comparison. On outlier workloads, uniform INT8 fails catastrophically ($\text{RMSE} > 1.0$), while the evolved policy matches FP16 exactly.




Figure 1: Log-scale relative RMSE across all 12 workloads. Evolved policy (green) tracks FP16 (gray) exactly. Uniform INT8 (orange) diverges by orders of magnitude on outlier workloads—RMSE > 1.0 means the output is noise.

Table 1: Average metrics across all 12 workloads.

Strategy	Avg. RMSE	Bits	Compression
Uniform FP16	4.93e-4	16.0	1.0×
Evolved Policy	1.08e-2	11.3	1.4×
Uniform INT8	5.61e-1	8.0	2.0×

4.3 Entropy as the Key Discriminator

The entropy distribution across all 19,488 evaluated blocks is strongly bimodal (Figure 2): 13,840 blocks (71.0%) cluster at entropy ≈ 4.3 (uniform attention, safe for INT8); 5,648 blocks (29.0%) cluster at entropy 0.5–1.0 (peaked attention from outlier workloads, assigned FP16). The gap between 1.5 and 3.5 is nearly empty, making the threshold robust to its exact value.

4.4 Sequence Length Invariance

The policy’s decisions are independent of sequence length (Figure 4). Across $N = 512$ to $N = 8,192$, compression is constant within each workload family, contradicting the hypothesis that longer sequences have proportionally more compressible blocks. The outlier/entropy signal dominates positional effects entirely.

4.5 Comparison to FlashAttention-2

We benchmark the evolved policy against PyTorch SDPA dispatching to FlashAttention-2 (FA-2) on an NVIDIA RTX A4000 (Ampere, sm.86). Table 3 reports per-workload relative RMSE against a FP32 tiled reference, alongside the direct output RMSE between policy and FA-2 (Match $< 5\%$).

All 12 workloads pass the match criterion (12/12). On the five outlier workloads, the policy achieves *lower* RMSE than FA-2: $0.37\text{--}0.73\times$ FA-2 error. This advantage arises because our policy accumulates in FP32 while FA-2 runs end-to-end in FP16—an effect most pronounced when outlier scores amplify FP16 rounding. On the seven non-outlier workloads the policy uses INT8 compression (8 bits, $2\times$ savings) and its RMSE rises to ≈ 0.018 , still within the match threshold. The KV-cache achieves a uniform $2\times$ compression (50% memory reduction) across all workloads and all sequence lengths from 512 to 8,192.

4.6 Entropy-to-Ratio Simplification

The evolved policy uses softmax entropy, which requires computing $\sum_i p_i \log p_i$ over every row of the attention block after a full softmax pass—approximately 32,000 transcendental operations per 128×128 tile. For ASIC implementation, transcendental operations are expensive. We show that a

Table 2: Per-workload relative RMSE vs. FP64 reference. Shaded rows are outlier workloads. Uniform INT8 is catastrophic there; evolved policy matches FP16.

N	d	Causal	Outliers	FP16 RMSE	INT8 RMSE	Evolved RMSE	Bits
512	64	No	No	1.84e-4	1.81e-2	1.81e-2	8
512	128	Yes	Yes	5.74e-4	1.33e+0	5.74e-4	16
2048	64	No	Yes	1.00e-3	1.35e+0	1.00e-3	16
2048	128	Yes	No	1.67e-4	1.79e-2	1.79e-2	8
8192	64	Yes	No	1.73e-4	1.62e-2	1.62e-2	8
8192	128	No	Yes	1.14e-3	1.38e+0	1.14e-3	16
2048	64	Yes	Yes	1.06e-3	1.27e+0	1.06e-3	16
2048	128	No	No	1.82e-4	1.92e-2	1.92e-2	8
4096	64	No	No	1.89e-4	1.92e-2	1.92e-2	8
4096	128	Yes	Yes	8.86e-4	1.28e+0	8.86e-4	16
8192	64	No	No	1.83e-4	1.75e-2	1.75e-2	8
8192	128	Yes	No	1.72e-4	1.68e-2	1.68e-2	8

Why It Works: Entropy Tells You Everything




Figure 2: Left: Bimodal block entropy histogram with the 2.0 threshold marked. Two clusters are clearly separated by a wide gap. Right: Precision decisions by entropy—clean separation at the threshold. INT8 blocks (blue dots) cluster at high entropy; FP16 blocks (red crosses) at low entropy.

Table 3: Evolved policy vs. FlashAttention-2 (PyTorch SDPA, RTX A4000). RMSE vs. FP32 reference. **Bold**: policy outperforms FA-2. All 12/12 workloads: direct output RMSE vs. FA-2 < 5%.

N	d	Outliers	FA-2 RMSE	Policy RMSE	Bits
512	64	No	4.6e-4	1.8e-2	8
512	128	Yes	1.2e-3	5.7e-4	16
2048	64	Yes	2.3e-3	1.0e-3	16
2048	128	No	4.2e-4	1.8e-2	8
8192	64	No	4.3e-4	1.6e-2	8
8192	128	Yes	2.3e-3	1.1e-3	16
2048	64	Yes	2.0e-3	1.1e-3	16
2048	128	No	4.6e-4	1.9e-2	8
4096	64	No	4.7e-4	1.9e-2	8
4096	128	Yes	1.6e-3	8.9e-4	16
8192	64	No	4.7e-4	1.8e-2	8
8192	128	No	4.3e-4	1.7e-2	8
Average			1.0e-3	1.1e-2	11.3




Figure 3: Per-block precision maps across all 12 workloads (blue=INT8, red=FP16). Each sub-plot is $\text{block_row} \times \text{block_col}$. Outlier workloads are uniformly FP16; non-outlier workloads are uniformly INT8. No mixed blocks appear within any single workload.




Figure 4: Compression vs. sequence length. Both lines are flat: the precision decision is length-invariant.

Table 4: Signal comparison: agreement with evolved policy on 19,488 blocks.

Signal	Agree	Danger miss
Entropy $\sum p \log p$	100%	0
Ratio $r > 10$	100%	0
Score variance	71%	many

pre-softmax ratio signal achieves identical precision decisions.

Ratio signal. Define the score matrix $\mathbf{S} = \mathbf{Q}\mathbf{K}^\top / \sqrt{d}$. The ratio signal is:

$$r = \frac{\max |\mathbf{S}|}{\text{mean} |\mathbf{S}|}$$

This measures how much one element dominates the block. For a near-uniform block, $r \approx 1$. For a block with a spike (high entropy gradient, few tokens attended), $r \gg 10$.

Benchmark. Across 19,488 blocks from our workload suite:

- FP16-assigned blocks: median $r = 957$ (outlier scores dominate)
- INT8-assigned blocks: median $r = 5.9$ (near-uniform scores)
- Gap between $r = 6$ and $r = 895$: zero overlap

At threshold $r = 10$: 100% agreement with the entropy-based policy, 0 dangerous misses (Table 4).

No-division hardware formulation. The threshold check $r > 10$ is equivalent to:

$$\max |\mathbf{S}| \times N > 10 \times \text{sum} |\mathbf{S}|$$

where $N = \text{BLOCK_M} \times \text{BLOCK_N}$ is a compile-time constant (power of two). This eliminates the divider entirely:

- $\max \times N$: left-shift by $\log_2 N$ (free—wire routing)
- $10 \times \text{sum}$: $(\text{sum} \ll 3) + (\text{sum} \ll 1)$ (two shifts and one adder)
- Result: one comparator, one output bit

The max and sum are already computed during the $\mathbf{Q} \cdot \mathbf{K}^\top$ scan (max is required for online softmax stability in FlashAttention), so the only new silicon is the multiply-by-10 and comparator.

5 Real LLM Validation

We validate the ratio signal on real language model activations using three members of the Qwen2/Qwen2.5 family and Microsoft Phi-2, spanning two distinct architecture families and model scales from 0.5B to 3B parameters.

5.1 Methodology

We register PyTorch forward hooks on the `q.proj`, `k.proj`, and `v.proj` linear layers of each attention block, capturing raw pre-softmax activations. This directly measures the quantity that the hardware controller observes. We handle grouped-query attention (GQA) by expanding KV heads via `repeat_interleave` before computing scores, and run 512-token sequences (the minimum length for meaningful tile statistics). Two prompts per model are used (prose and code), giving consistent coverage across different attention pattern types. Phi-2 uses standard multi-head attention (MHA, 32 Q and 32 KV heads, `head_dim=80`), making it a structurally distinct datapoint from the GQA-based Qwen2 models.

An earlier attempt measuring post-softmax entropy on 20–54-token sequences yielded 89.7% FP16-flagged tiles—an artefact of the wrong metric (post-softmax probabilities always concentrate on short causal sequences) and insufficient context length. All results below use the correct pre-softmax ratio on 512-token sequences.

5.2 Per-Model Results

Table 5 reports INT8-safe tile percentage, median ratio, and maximum observed ratio across 3,000–4,600 tiles per model.

Table 5: Real LLM validation: pre-softmax ratio at 512 tokens. INT8-safe = ratio ≤ 10 . Phi-2 uses MHA; Qwen2 models use GQA.

Model	Attn	L	INT8	Med. r	Max r
Qwen2-0.5B	GQA	24	91.7%	5.31	16.2
Qwen2-1.5B	GQA	28	97.1%	4.20	14.3
Qwen2.5-3B	GQA	36	96.6%	4.07	32.5
Phi-2 (2.7B)	MHA	32	58.9%	5.93	41.0

5.3 Cross-Architecture Analysis

Qwen2 family. INT8 coverage *improves* from 91.7% at 0.5B to 97.1% at 1.5B, then holds at 96.6% at 3B. The median ratio decreases monotonically (5.31 $\rightarrow$ 4.07), indicating score distributions become more uniform at scale. FP16 tiles consistently concentrate at layer 0 (chaotic attention before the residual stream stabilises) and at one or two mid-network layers. Layers 1 through $L-2$ are 0% FP16 in both the 1.5B and 3B models. KV channel outliers remain $\leq 0.8\%$ per layer.

Phi-2 (2.7B). Phi-2 exhibits a qualitatively different pattern: 41.1% FP16 tiles, with a bimodal per-layer distribution where individual layers are either 0% or 100% FP16. This reflects Phi-2’s training objective and architectural choices (dense MHA, no GQA, wider head dimension of 80 vs. 64). Critically, the ratio gate behaves *correctly*—it identifies the high-ratio layers and routes them to FP16, protecting numerical accuracy. The signal is honest, not optimistically biased: when a model genuinely requires FP16 precision, the controller says so. KV channel outliers are higher in Phi-2’s early layers (up to 32.8% at layer 29) but near-zero in the majority of layers.

The results confirm that the ratio gate is architecture-agnostic: it measures a physical property of the attention score distribution and makes the correct decision regardless of model family or attention variant.

5.4 Model Weight Streaming

To avoid exhausting disk storage (4.1 GB free at time of experiments), model weights are streamed through `/dev/shm` (Linux RAM-backed tmpfs, 11 GB available). The workflow per model is: download shards to `/dev/shm` $\rightarrow$ load weights to GPU VRAM $\rightarrow$ delete `/dev/shm` cache (weights now reside only in VRAM) $\rightarrow$ run forward hooks $\rightarrow$ delete GPU model and call `torch.cuda.empty_cache()`. Peak `/dev/shm` usage is 6.2 GB (Qwen2.5-3B shards); disk usage is zero.

Table 6: Fixed-point simulation: agreement with float reference. A dangerous miss is INT8 called when float says FP16 (catastrophic: RMSE 0.001 $\rightarrow$ 1.5). 0 dangerous misses at all tested widths.

Bits	Agreement	Danger misses	False pos.
4	100.0%	0	0
6	100.0%	0	0
8	100.0%	0	0
10	100.0%	0	0
12	100.0%	0	0
16	100.0%	0	0

6 Fixed-Point Hardware Simulation

We simulate the precision controller in fixed-point integer arithmetic to verify that the ratio decision is robust to quantization of the scores themselves. This answers the chip design question: at what register bit width does the decision become incorrect?

6.1 Formulation

For a tile of scores quantized to b -bit integers with per-tile max-abs scaling:

1. $\mathbf{S}_{\text{int}} = \text{round}(\mathbf{S}/s)$, $s = \max |\mathbf{S}|/(2^{b-1} - 1)$
2. $m = \max |\mathbf{S}_{\text{int}}|$ (b -bit register)
3. $\sigma = \sum |\mathbf{S}_{\text{int}}|$ ($(b+12)$ -bit register for 64×64 tiles)
4. Decision: $m \times 4096 > 10 \times \sigma$ (i.e., $m \ll 12$ vs. $(\sigma \ll 3) + (\sigma \ll 1)$)

6.2 Results

Table 6 reports agreement with the float reference over 8,000 synthetic tiles (7,360 INT8-ground-truth, 640 FP16-ground-truth) at bit widths from 4 to 16.

Even 4-bit score representation gives perfect agreement. The ratio between FP16 and INT8 tile medians (957 vs. 5.9) is so large that quantization noise cannot bridge the gap. The minimum-hardware register budget is therefore:

Register	Width
Max accumulator	4-bit (minimum proven safe)
Sum accumulator	16-bit ($4 + \log_2 4096$)
LHS ($m \ll 12$)	16-bit (free: wire routing)
RHS ($10 \times \sigma$)	20-bit (shift-add)
Comparator	20-bit, 1-bit output

7 KV-Cache Application

We implement an entropy-guided KV-cache quantization module with four hardware-mapped components:

1. **Precision Controller.** Ratio $> 10 \rightarrow$ 1-bit precision select. Hardware: shift-add and comparator.
2. **Quantizer.** Symmetric INT8 (max-abs scan, multiply-round-clamp) or FP16 passthrough. Hardware: multiplier pipeline with precision-select mux.
3. **Cache SRAM.** Fixed header: [1-bit tag] [16-bit scale] [$N \times D$ payload].
4. **Dequantizer.** Reads tag, applies inverse. Hardware: single multiplier pipeline.

Benchmarks across all 12 workloads confirm a uniform $2\times$ compression ratio (50% memory reduction) at 8 bits/element, scaling flat from $N = 512$ to $N = 8,192$. This is validated by 45 passing unit tests covering correctness (MAE < 0.02), autoregressive append, and multi-layer storage.

8 Triton Kernel Prototype

We implement the precision policy as a Triton kernel on an NVIDIA RTX A4000. The kernel follows the FA-2 tile loop exactly, adding three lines after computing $S = QK^\top$:

```
abs_s = tl.abs(s)
s_max = tl.max(abs_s)
s_mean = tl.sum(abs_s) / (BLOCK_M *
    BLOCK_N)
use_int8 = (s_max / (s_mean + 1e-6)) <
    RATIO_THRESHOLD
```

Listing 2: Precision decision inside the Triton tile loop (software).

In hardware the division is replaced by the no-division formulation from Section 4.6: $\max \ll \log_2 N > (\sigma \ll 3) + (\sigma \ll 1)$.

When `use_int8` is true, K and V are symmetrically quantized to INT8 and dequantized before the dot products; otherwise FP16 values are used unchanged. All 12 workloads pass correctness tests (kernel output within $3\times$ SDPA error vs. FP32 reference), and all 12/12 precision decisions match the ratio > 10 rule.

The kernel currently runs $2.3\text{--}8.7\times$ slower than PyTorch SDPA because Triton’s `tl.dot` always uses

FP16 compute—simulated INT8 adds quantize/dequantize roundtrips rather than saving multiplier cycles. The throughput gain materialises only on hardware with integer multiply units (FPGA DSPs in INT8 mode, or ASIC integer multipliers). Block sizes are 128×128 for $d = 64$ and 64×64 for $d = 128$, constrained by the A4000’s 101,376-byte per-SM shared memory limit.

9 Hardware Design Implications

9.1 Why Simple Beats Learned

The search had full freedom to evolve arbitrary logic over all 14 features. It converged on two comparisons—evidence that the problem is genuinely simple. A minimal MLP controller would require multipliers, weight SRAM, and activation logic on the critical path. The threshold comparator requires a max tree, an adder tree, and one comparator. The fixed-point simulation confirms that 4-bit score registers suffice.

9.2 Programmable Threshold Registers

We propose storing per-layer thresholds in a small register file:

```
precision = (ratio > THRESH[1]) ? FP16 :
    INT8;
```

where `ratio` is the integer comparison result (1 bit) and `THRESH[1]` can be tuned per layer. For a 32-layer transformer this costs 64 bytes. Real LLM validation (Section 5) shows that layer 0 consistently benefits from a higher threshold or a hardcoded FP16 flag; layers $1\text{--}L-2$ are fully INT8-safe and do not require threshold tuning. Programmed once at model load time from offline profiling, this provides per-layer adaptation with no additional logic—capturing $\sim 97\%$ of the benefit of a learned controller at $\sim 0\%$ of the silicon cost.

9.3 Gate-Count Summary

The complete precision controller for a 64×64 tile:

Component	Implementation
Max tree	12 comparator levels
Adder tree	12 adder levels
LHS shift	Wire routing (free)
RHS $\times 10$	2 shifts + 1 adder
Comparator	1 comparator, 1-bit output
Threshold regs	64 bytes (32 layers)

The max tree and adder tree are already required by the FlashAttention online softmax (max for numerical stability, sum for normalisation). The only new logic is the multiply-by-10 and the comparator.

10 RTL Verification and Synthesis

The gate-count estimate above is verified end-to-end in three stages: a self-checking SystemVerilog testbench, a real-data replay testbench, and Liberty-mapped logic synthesis. All three flows are reproducible from `rtl/Makefile`.

10.1 Functional Verification

The DUT is `rtl/precision_controller.sv`, a streaming controller that consumes one signed score per cycle, asserts `s_last` on the final score of each tile, and produces a registered FP16/INT8 decision exactly one cycle later. Two testbenches drive the same DUT:

- **Directed TB** (`tb_precision_controller.sv`): 110 synthetic tiles covering uniform constants, single spikes at boundary positions, alternating signs, and 100 LCG-random tiles. Reference computed in software using the exact integer formula $\max N > T \cdot \Sigma$. Result: 110/110 pass.
- **Replay TB** (`tb_realdata.sv`): 143 tiles generated by `analysis/gen_rtl_testvectors.py`, including INT8-quantized samples from real attention-score distributions and explicit decision-edge cases. The two canonical edges: with background = 1 and one spike, spike = 10 gives LHS = 40960 < RHS = 41050 (INT8), and spike = 11 gives LHS = 45056 > RHS = 41060 (FP16). Tiles are streamed via `$readmemh`. Result: 143/143 pass.

The replay tiles are produced from the same INT8 quantization the fixed-point simulation validated, so passing the replay TB demonstrates bit-exact agreement between Python reference and SystemVerilog DUT on realistic data, including both sides of the decision threshold.

10.2 Parameter Sweep

We synthesized the DUT in Yosys with `abc -fast -g NAND` mapping across two axes: tile size $N = \text{BLOCK_M} \times \text{BLOCK_N}$ from 256 to 65,536 at `SCORE_WIDTH = 8`, and bit-width from 4 to 16 at fixed tile 64×64 . The flip-flop count matches the closed-form expression at every point, with no inferred state:

$$\text{FFs} = 2 \cdot \text{SCORE_WIDTH} + \log_2 N + 2$$

Tile area is logarithmic, not linear: a $256 \times$ increase in tile size ($16 \times 16 \rightarrow 256 \times 256$) adds only +34% NAND2-equivalent gates ($737 \rightarrow 988$) and +8 flip-flops. Bit-width is the dominant cost: at fixed tile 64×64 , going from 4 to 16 bits more than doubles area ($588 \rightarrow 1431$ NAND2-eq) and FFs ($22 \rightarrow 46$). This confirms the fixed-point sim’s choice of 8 bits as the sweet spot.

10.3 Liberty-Mapped Synthesis (ASAP7)

For technology-aware area and timing we synthesized against ASAP7 [Clark et al., 2016], an open 7 nm FinFET enablement that serves as the closest public proxy for our intended TSMC 16FFC target. Yosys 0.9 with `dfflibmap + abc -liberty` on the RVT FF corner produced (deterministic across reruns):

Tile	FFs	Cell area (μm^2)	Critical path (ps)
64×64	30	28.24	417.74
128×128	32	31.39	465.74

The reference 64×64 controller is 30 FFs and just under $30 \mu\text{m}^2$ in ASAP7—genuinely negligible silicon. The longest combinational path is the `sum_acc` adder $\rightarrow$ shift-add $\rightarrow$ 24-bit comparator $\rightarrow$ `d_fp16` register at 417.74 ps.

10.4 Sky130 OpenLane End-to-End Sign-off

To complement the ASAP7 area/timing estimate with a real implementation, we re-synthesized the controller through the LibreLane / OpenROAD [LibreLane Contributors, 2024] open-source ASIC flow targeting SkyWater Sky130A. Unlike the Yosys+abc pass, this is a complete industrial-quality flow: synthesis, floorplanning, placement, clock-tree synthesis, global and detailed routing, parasitic extraction, post-PnR multi-corner static timing analysis, DRC, LVS, antenna, and IR-drop. Total runtime is ~ 3 minutes

on a 4-core machine. The signed-off layout is shown in Figure 5.




Figure 5: Sky130A precision controller layout, signed off at 80 MHz. Die size $87.8 \times 98.5 \mu\text{m}^2$ at 59% core utilization; 30 flip-flops and 292 logic cells; total routed wirelength 6,201 μm .

Sign-off was performed across the three Sky130A voltage/temperature corners (Table 7). The SS corner is the gating condition; the design closes with +72 ps margin and zero setup, hold, DRC, LVS, antenna, or IR-drop violations.

Table 7: Sky130A multi-corner sign-off at 80 MHz (12.5 ns period). Zero violations of any kind across all corners.

Corner	Setup WNS	Hold WS
FF (-40°C , 1.95 V)	+8.342 ns	+0.182 ns
TT ($+25^\circ\text{C}$, 1.80 V)	+6.102 ns	+0.363 ns
SS ($+100^\circ\text{C}$, 1.60 V, sign-off)	+0.072 ns	+0.860 ns

Headline implementation numbers at Sky130A, 80 MHz: 3,438 μm^2 stdcell area in 8,642 μm^2 die; 330.5 μW total power at the TT corner (91.5 μW switching, 238.9 μW internal, 6 pW leakage); 30 flip-flops, exactly matching the closed-form $2 \cdot \text{SCORE_WIDTH} + \log_2 N + 2$ prediction in our gate-count summary.

Tuning note. LibreLane’s first-pass defaults (IO_DELAY_CONSTRAINT = 20%, CLOCK_UNCERTAINTY = 250 ps) are chip-level conservative and produced negative SS slack of -1.58 ns. Three knobs landed the clean sign-off: IO_DELAY_CONSTRAINT = 5% (appropriate for a small block embedded in a larger pipeline rather than a top-level chip), CLOCK_UNCERTAINTY = 100 ps (matching the sub-50 ps CTS skew of a design this size), and DESIGN_REPAIR_MAX_SLEW_PCT = 0 (force the repair tool to fix every slew violation, not just

the worst 20%). The worst-case timing path was input-to-flip-flop, not flip-flop-to-flip-flop, which is why pure clock-period increases gave diminishing returns until the I/O budget was tightened.

Cross-check with the ASAP7 projection. Sky130 to 16FFC scales by approximately $5\times$ area reduction, $2\times$ speed gain at TT, and $10\times$ dynamic power reduction (NAND2 and DFF ratios from published process comparisons). Applied to the Sky130 signed-off numbers, this projects to $\sim 700 \mu\text{m}^2$, ~ 600 MHz, and $\sim 30 \mu\text{W}$ in TSMC 16FFC—consistent with the ASAP7-derived projection in Section 10.5 and with the 800 MHz target encoded in our Cadence SDC constraints. Two independent industrial-grade flows now agree on the design’s silicon characteristics.

10.5 TSMC 16FFC Projection

Scaling ASAP7 results to TSMC 16FFC using published per-cell area and delay ratios (NAND2: $\sim 5\times$, DFF: $\sim 4\times$; transistor delay: $\sim 2\times$ at TT, $\sim 3\times$ at SS):

Corner	Critical path (ps)	Slack at 800 MHz
TSMC 16FFC FF (best)	~ 580	+670
TSMC 16FFC TT (typical)	~ 835	+410
TSMC 16FFC SS (worst, sign-off)	~ 1255	≈ 0

Projected cell area at the reference 64×64 , 8-bit configuration is $\sim 115\text{--}160 \mu\text{m}^2$ —roughly 10–20% of a single FP16 multiplier and well under 0.2% of a typical 8×8 INT8 MAC array in 16FFC. The decision overhead is rounding error against the compute it gates.

The SS corner has effectively zero slack at 800 MHz. If full sign-off in Genus on the real PDK confirms negative slack, we have a documented single-register pipeline patch (one stage between `sum_next` and the comparator) that halves the critical path, adds 24 flip-flops, and delays the decision by one cycle—a non-issue since the upstream attention pipeline is much deeper.

The full Cadence Genus and Innovus scripts (`rtl/genus.tcl`, `rtl/innovus.tcl`, `rtl/mmmc.tcl`) are parameterized over the PDK paths and ready for the TSMC University Program shuttle.

11 Limitations

- (1) **Workload-level granularity**—no fine-grained per-block decisions within outlier sequences.
- (2) **Fixed outlier threshold**—not co-evolved

with the ratio threshold, may require re-tuning for architectures with very different score magnitudes. (3) **No throughput measurement**—the Triton kernel proves correctness but not speedup; that requires hardware integer multipliers. (4) **Dense attention only**—grouped-query, sliding window, and sparse patterns require separate validation (though GQA is handled correctly in our real LLM validation via head expansion). (5) **Three model families tested**—Qwen2, Qwen2.5, and Phi-2; Llama, Mistral, and Gemma remain as future validation targets.

12 Future Work

(1) **Additional architectures.** Validate the ratio signal on Llama-3, Mistral, and Gemma. Phi-2 results (58.9% INT8-safe) suggest the threshold may benefit from per-architecture tuning; a learnable or workload-calibrated threshold is a natural extension. (2) **Sign-off on TSMC 16FFC.** ASAP7 synthesis (Section 10.3) and Sky130 end-to-end PnR (Section 10.4) independently project to $\sim 150 \mu\text{m}^2$ and ~ 0 ps slack at the SS corner @ 800 MHz on TSMC 16FFC. Real Genus and Innovus runs on the University Program shuttle will confirm the projection and yield post-PnR power numbers from the replay-TB activity dump. If SS slack is negative, the documented single-register pipeline patch closes timing. (3) **End-to-end throughput.** Integrate the controller into a full attention datapath to measure real speedup from INT8 tile routing in hardware. (4) **Co-evolution of thresholds.** Jointly evolve the ratio threshold and per-layer register values across model families. (5) **H100 comparison.** Test against FlashAttention-3 FP8 on Hopper when hardware is available.

13 Conclusion

We have traced a complete path from evolutionary policy discovery to hardware-verified implementation. The search converges on softmax entropy as the key discriminator; we then show entropy is replaceable by a pre-softmax ratio requiring only a max tree, an adder tree, and a comparator—no transcendental operations. Fixed-point simulation confirms the decision is correct at 4-bit score representation with zero dangerous misses across 8,000 tiles. Real LLM validation on Qwen2/Qwen2.5 models (0.5B–3B, 24–36 layers) shows 91–97% of attention tiles are INT8-safe, with the fraction *improving* at larger scale. The SystemVerilog implementation

passes 253 directed and real-data testbench cases bit-exactly, synthesizes to 30 flip-flops and $\sim 28 \mu\text{m}^2$ in ASAP7 (open 7 nm proxy), runs end-to-end through the LibreLane/OpenROAD industrial open flow on Sky130 with a fully signed-off layout at 80 MHz (zero DRC/LVS/antenna/IR-drop/ setup/hold violations across all three corners), and projects to $\sim 150 \mu\text{m}^2$ with +415 ps slack at the typical TSMC 16FFC corner @ 800 MHz—a precision controller ready for the University Program shuttle.

References

- Zahra V. Boroujeni et al. Adaptive KV-cache quantization for long-context inference. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2026.
- Lawrence T. Clark, Vinay Vashishtha, Lucian Shifren, Aditya Gujja, Saurabh Sinha, Brian Cline, Chandrasekaran Ramamurthy, and Greg Yeric. ASAP7: A 7-nm finfet predictive process design kit. *Microelectronics Journal*, 53:105–115, 2016.
- CodeLion. OpenEvolve: Evolutionary code optimization framework. <https://github.com/codelion/openevolve>, 2025.
- Tri Dao. FlashAttention-2: Faster attention with better parallelism and work partitioning. In *International Conference on Learning Representations*, 2024.
- Tri Dao, Dan Fu, Stefano Ermon, Atri Rudra, and Christopher Ré. FlashAttention: Fast and memory-efficient exact attention with IO-awareness. In *Advances in Neural Information Processing Systems*, volume 35, pages 16344–16359, 2022.
- Elias Frantar, Saleh Ashkboos, Torsten Hoefer, and Dan Alistarh. GPTQ: Accurate post-training quantization for generative pre-trained transformers. In *International Conference on Learning Representations*, 2023.
- LibreLane Contributors. LibreLane: An open-source ASIC implementation flow built on OpenROAD. <https://github.com/librelane/librelane>, 2024.
- Bernardino Romera-Paredes et al. Mathematical discoveries from program search with large language models. *Nature*, 625:468–475, 2024.
- Jay Shah, Ganesh Bikshandi, Ying Zhang, Vijay Thakkar, Pradeep Ramani, and Tri Dao. FlashAttention-3: Fast and accurate attention

with asynchrony and low-precision. *arXiv preprint arXiv:2407.08608*, 2024.

Guangxuan Xiao, Ji Lin, Mickael Seznec, Hao Wu, Julien Demouth, and Song Han. SmoothQuant: Accurate and efficient post-training quantization for large language models. In *International Conference on Machine Learning*, 2023.